

Three Dark Clouds over the Android Kernel

Jun Yao @ Alpha Lab

About me



- **Security researcher in Alpha Lab of 360**
- **Focus on the Android/Linux kernel**
 - ✓ Bug hunting
 - ✓ Writing exploits
 - ✓ Researching Mitigations

Contents



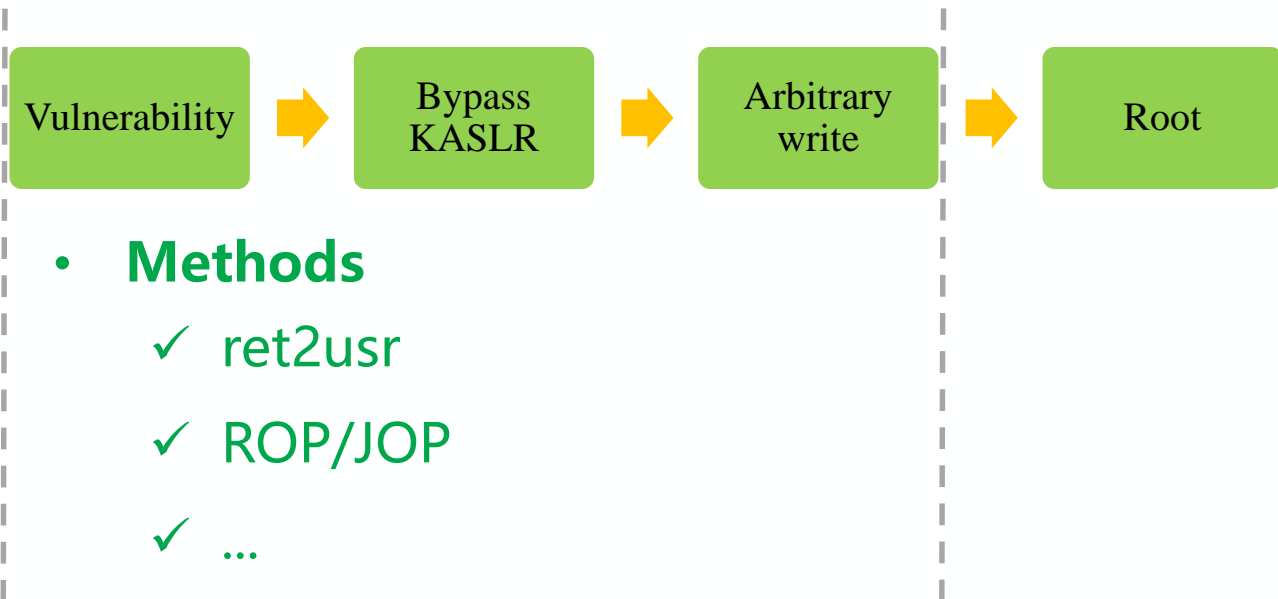
- 01.** Three Dark Clouds
- 02.** Thunder and Lightning
- 03.** Dispel the Clouds and See the Sun
- 04.** Conclusion

Contents




- 01.** Three Dark Clouds
- 02.** Thunder and Lightning
- 03.** Dispel the Clouds and See the Sun
- 04.** Conclusion

Classic attack path




The ret2dir



The ret2dir

- **Double Mapping**



<https://cs.brown.edu/~vpk/papers/ret2dir.sec14.pdf>

The ret2dir



- **User can calculate the kernel address**
 1. Get PFN from `/proc/self/pagemap`
 2. $kaddr = (PFN \ll PAGE_SHIFT) - PHYS_OFFSET + PAGE_OFFSET$
- **The kernel mapping attributes**
 - ✓ qcom: RW
 - ✓ mtk: RWX (CVE-2016-0820)

- **Hide PFN in the /proc/self/pagemap**

```
+static int pagemap_open(struct inode *inode, struct file *file)
+{
+    /* do not disclose physical addresses: attack vector */
+    if (!capable(CAP_SYS_ADMIN))
+        return -EPERM;
+    return 0;
+}
+
+const struct file_operations proc_pagemap_operations = {
+    .llseek    = mem_llseek, /* borrow this */
+    .read      = pagemap_read,
+    .open      = pagemap_open,
+};
```

<https://android.googlesource.com/kernel/msm/+718c232053ec81a4a732ed50ba0224ae512bb694%5E%21/>

The ret2dir




- **It still works**
 - ✓ CVE-2015-3636^[1]
 - ✓ CVE-2018-9568^[2]
- **The key is information leakage**

[1] <https://www.blackhat.com/docs/us-15/materials/us-15-Xu-Ah-Universal-Android-Rooting-Is-Back.pdf>

[2] <https://github.com/ThomasKing2014/slides/blob/master/Building%20universal%20Android%20rooting%20with%20a%20type%20confusion%20vulnerability.pdf>

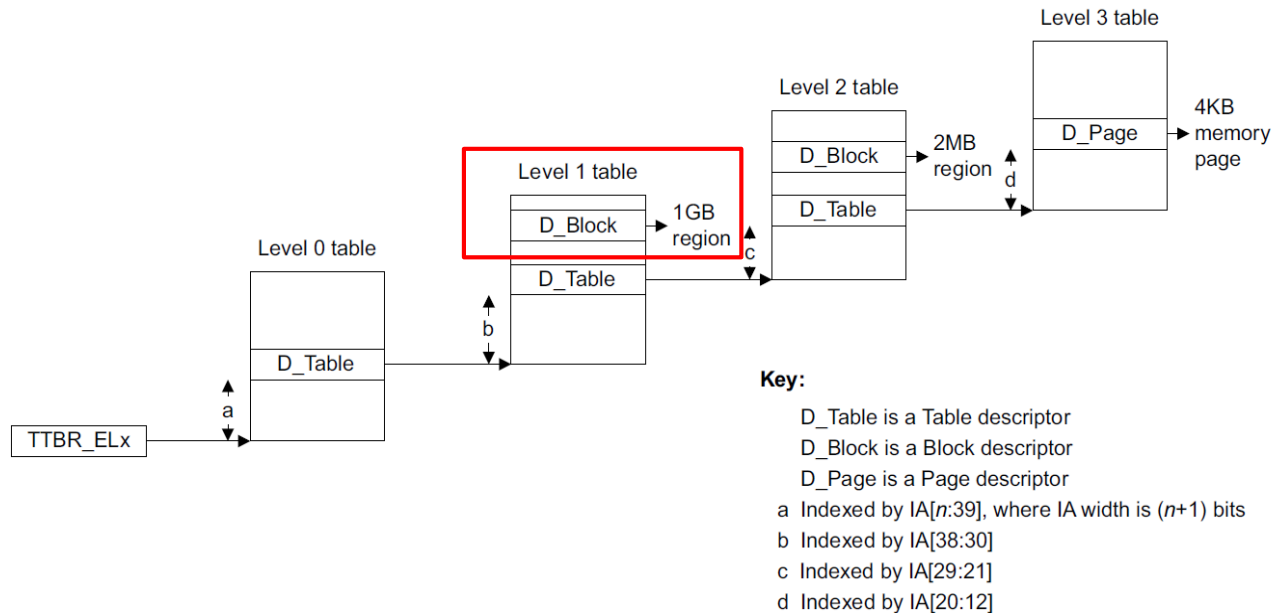
- **Kernel Space Mirror Attack**



<https://i.blackhat.com/briefings/asia/2018/asia-18-WANG-KSMA-Breaking-Android-kernel-isolation-and-Rooting-with-ARM-MMU-features.pdf>


The KSMA

- The block entry



<https://documentation-service.arm.com/static/5f20515cbb903e39c84dc459?token=>

- The AP (Access Permissions) attribute

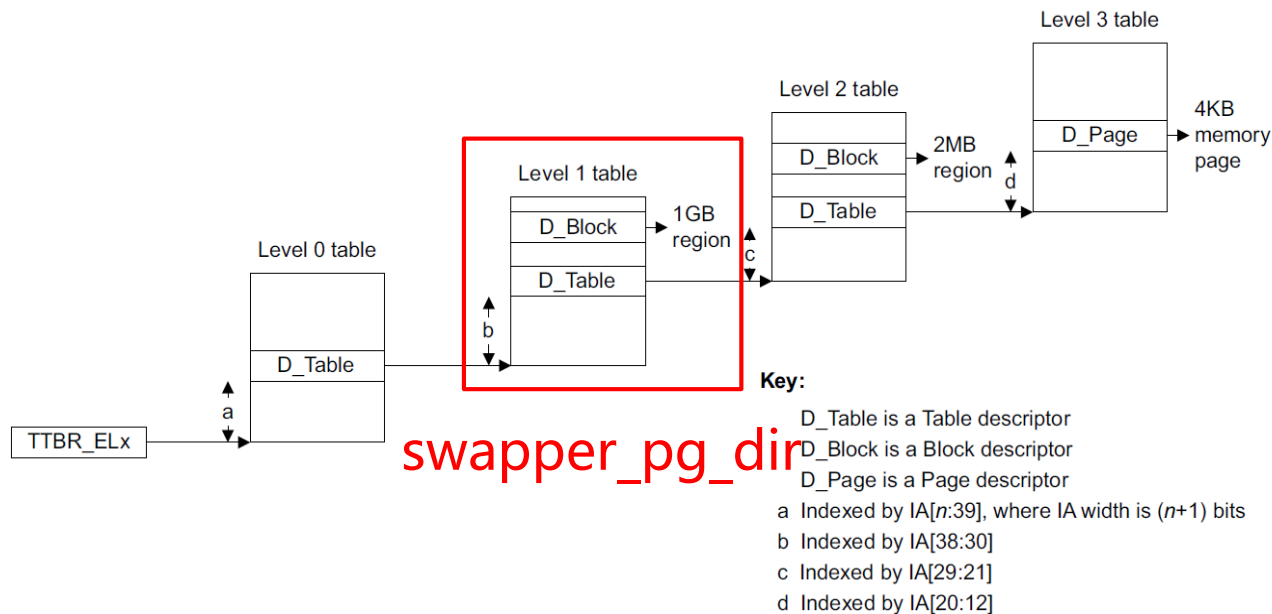


AP[2:1]	Access from higher Exception level	Access from EL0
00	Read/write	None
01	Read/write	Read/write
10	Read-only	None
11	Read-only	Read-only

The KSMA



- Android: 39-bits VA & 4K page



- **The position of the swapper_pg_dir is fixed**

arch/arm64/kernel/vmlinux.lds.S

```
BSS_SECTION(0, 0, 0)
```

```
. = ALIGN(PAGE_SIZE);  
idmap_pg_dir = .;  
. += IDMAP_DIR_SIZE;  
swapper_pg_dir = .;  
. += SWAPPER_DIR_SIZE;
```

- **Arbitrary write is needed**

The PIPE



- **Proposed in 2017**
- **Has the characteristics of TOCTOU**
- **Converts heap problems to arbitrary reading and writing**

- Allocate iovec on the heap

fs/read_write.c

```
ssize_t rw_copy_check_uvector(int type, const struct iovec __user * uvector,  
                               unsigned long nr_segs, unsigned long fast_segs,  
                               struct iovec *fast_pointer, ...) {  
    [...]   
    if (nr_segs > fast_segs) { // UIO_FASTIOV == 8  
        iov = kmalloc(nr_segs*sizeof(struct iovec), GFP_KERNEL);  
        [...]   
    }  
    if (copy_from_user(iov, uvector, nr_segs*sizeof(*uvector))) {  
        ret = -EFAULT;  
        goto out;  
    }  
}
```

- Check the iovec passed in by the user

fs/read_write.c

```
ssize_t rw_copy_check_uvector(int type, const struct iovec __user * uvector,
                               unsigned long nr_segs, unsigned long fast_segs,
                               struct iovec *fast_pointer, ...) {
    [...]
    for (seg = 0; seg < nr_segs; seg++) {
        void __user *buf = iov[seg].iov_base;
        ssize_t len = (ssize_t)iov[seg].iov_len;

        if (type >= 0 && unlikely(!access_ok(vrfy_dir(type), buf, len))) {}
        ret += len;
    }
}
```

The PIPE

- **We can block the pipe**
 1. There is no contents when reading
 2. There is no space when writing
- **Overwrite the iovec**



- Does **NOT** check when copying from/to the user

lib/iov_iter.c

```
static size_t copy_page_to_iter_iovec(..) {  
    [...]  
    left = __copy_to_user_inatomic(buf, from, copy);  
    copy -= left;  
    skip += copy;  
    from += copy;  
    bytes -= copy;  
    [...]  
}
```

- **Two mitigations**
 1. UAO (User Access Override)
 2. `uaccess_mask_ptr()`

`arch/arm64/include/asm/uaccess.h`

```
static inline void __user *__uaccess_mask_ptr(const void __user *ptr)
{
    asm volatile(
        "    bics xzr, %1, %2\n"
        "    csel %0, %1, xzr, eq\n"
        : "=&r" (safe_ptr)
        : "r" (ptr), "r" (current_thread_info()->addr_limit)
        : "cc");
    return safe_ptr;
}
```

<https://developer.arm.com/docs/ddi0595/h/aarch64-system-registers/ua0>

The SMA



- **SLAB Mirror Attack**
 - ✓ Two objects share the same SLAB
 - ✓ Bypass KASLR
 - ✓ Write kernel arbitrary (+KSMA = ROOT)

Contents

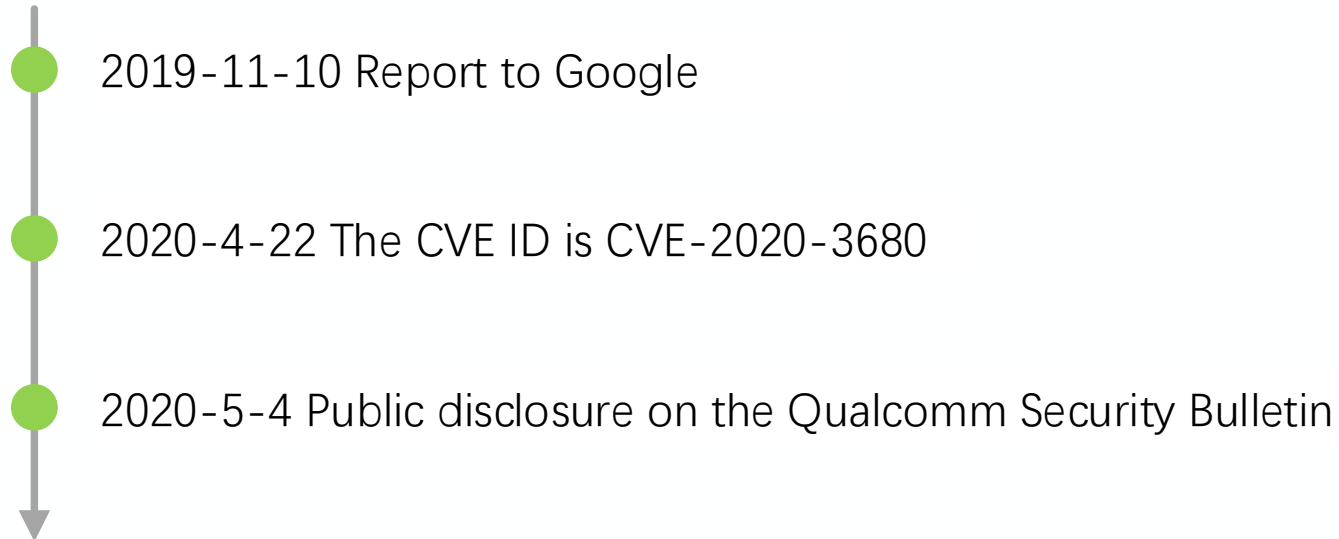


01. Three Dark Clouds
02. Thunder and Lightning
03. Dispel the Clouds and See the Sun
04. Conclusion

CVE-2020-3680



- **Timeline**





- Qualcomm ADSPRPC driver

drivers/char/adsrpc.c

```
int fastrpc_internal_mmap(...) {
    mutex_lock(&fl->map_mutex);
    [...]
    mutex_lock(&fl->fl_map_mutex);
    VERIFY(err, !fastrpc_mmap_create(..., &map));
    mutex_unlock(&fl->fl_map_mutex); // 1
    VERIFY(err, !fastrpc_mmap_on_dsp());
bail:
    if (err && map) {
        mutex_lock(&fl->fl_map_mutex);
        fastrpc_mmap_free(map, 0); // 4
        mutex_unlock(&fl->fl_map_mutex);
    }
    mutex_unlock(&fl->map_mutex);
    return err;
}
```

```
int fastrpc_internal_munmap_fd(...) {
    mutex_lock(&fl->fl_map_mutex); // 2
    if (fastrpc_mmap_find(..., &map))
        [...]
    if (map)
        fastrpc_mmap_free(map, 0); // 3
    mutex_unlock(&fl->fl_map_mutex);
    return err;
}
```

How to exploit



- **ROP/JOP** 🤔
 - ✓ PAN (Privileged Access Never)
 - ✓ CFI (Control Flow Integrity)
- **SMA + KSMA** 😊
 1. Convert UAF to Double Free
 2. Bypass KASLR
 3. Apply KSMA

Convert UAF to Double Free

drivers/char/adsrpc.c

```
void fastrpc_mmap_free(struct fastrpc_mmap *map, uint32_t flags) {  
    [...]  
    if (map->flags == ADSP_MMAP_HEAP_ADDR || ...) {  
        [...]  
    } else {  
        map->refs--;  
        if (!map->refs)  
            hlist_del_init(&map->hn); // 1  
        if (map->refs > 0 && !flags)  
            return;  
    }  
    if (map->flags == ADSP_MMAP_HEAP_ADDR || ...) {  
        [...]  
    } else if (map->flags == FASTRPC_DMAHANDLE_NOMAP) {  
        if (!IS_ERR_OR_NULL(map->handle))  
            ion_free(fl->apps->client, map->handle); // 2  
    }  
    kfree(map); // 3  
}
```

Convert UAF to Double Free



- How to bypass `hlist_del_init()`

```
include/linux/list.h
```

```
static inline void hlist_del_init(struct hlist_node *n) {  
    if (!hlist_unhashed(n)) {  
        __hlist_del(n);  
        INIT_HLIST_NODE(n);  
    }  
}
```

```
static inline int hlist_unhashed(const struct hlist_node *h  
{  
    return !h->pprev;  
}
```

Convert UAF to Double Free



- **The object used for heap spray**
 1. It uses the kmalloc-256 SLAB
 2. I can control the map->refs and map->flags
 3. The map->hn->pprev is zero

Convert UAF to Double Free



- The object used for heap spray


fs/xattr.c

```
setxattr(..., const void __user *value, size_t size, int flags) {  
    [...]  
    if (size) {  
        kvalue = kmalloc(size, GFP_KERNEL | __GFP_NOWARN);  
        if (copy_from_user(kvalue, value, size)) {  
            error = -EFAULT;  
            goto out;  
        }  
    }  
}
```

Convert UAF to Double Free



- The SLAB appears twice in the freelist



Bypass KASLR

- Use the freelist to modify the object

```
struct seq_file {  
    char *buf; // is modified by the freelist  
    size_t size;  
    [...]  
    struct mutex lock;  
    const struct seq_operations *op; // is leaked  
    int poll_event;  
    const struct file *file;  
    void *private;  
}
```

Bypass KASLR


- Users can read the seq_file->buf

```
flame:/ $ cat /proc/cpuinfo
Processor       : AArch64 Processor rev 14 (aarch64)
processor       : 0
BogoMIPS        : 38.00
Features        : fp asimd evtstrm aes pmull sha1 sha2 crc32 atomics
CPU implementer : 0x51
CPU architecture: 8
CPU variant     : 0xd
CPU part        : 0x805
CPU revision    : 14
```

Bypass KASLR




- Allocate victim_fd0 & victim_fd1




Bypass KASLR



- **Allocate the leaked_fd**




- **Free the leaked_fd**



Bypass KASLR



- Free the victim_fd0



- Read seq_file->op from victim_fd1

Apply KSMA



- **Use the object to modify the freelist**
 1. It uses the kmalloc-256 SLAB
 2. I can control the first 8 bytes

- Use the object to modify the freelist

```
struct ipv6_mc_socklist {  
    struct in6_addr addr;  
    int ifindex;  
    struct ipv6_mc_socklist __rcu *next;  
    rwlock_t sflock;  
    unsigned int sfmode;  
    struct ip6_sf_socklist *sflist;  
    struct rcu_head rcu;  
}
```


- Use the object to modify the freelist

net/ipv6/mcast.c

```
int ipv6_sock_mc_join(..., const struct in6_addr *addr) {  
    [...]  
    mc_lst = sock_kmalloc(sizeof(struct ipv6_mc_socklist));  
    if (!mc_lst)  
        return -ENOMEM;  
    mc_lst->next = NULL;  
    mc_lst->addr = *addr;  
    [...]  
}
```


Apply KSMA

- Allocate the `mc_lst`



Apply KSMA

- Update the `mc_lst->addr`



Demo



```
sirius:/data/local/tmp $
```

Contents



01. Three Dark Clouds
02. Thunder and Lightning
03. Dispel the Clouds and See the Sun
04. Conclusion

The ret2dir mitigation



- **Exclusive Page-Frame Ownership**
 - ✓ Unmap the user page in the kernel
 - ✓ Not yet merged into the mainline

<https://lwn.net/Articles/784839/>

The KSMA mitigation



- The KSMA patch
 - ✓ I submitted it in May 2018
 - ✓ Was merged into the mainline in September 2018
 - ✓ Android has not yet been backported 🙄

<https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/log/?h=v5.8.8&qt=author&q=yaojun>

The KSMA patch



- **Move the page table to the rodata section**

arch/arm64/kernel/vmlinux.lids.S

```
+#define KERNEL_PG_TABLES \  
+ . = ALIGN(PAGE_SIZE); \  
+ idmap_pg_dir = .; \  
+ . += IDMAP_DIR_SIZE; \  
+ TRAMP_PG_TABLE \  
+ RESERVED_PG_TABLE \  
+ swapper_pg_dir = .; \  
+ . += PAGE_SIZE; \  
+ swapper_pg_end = .;
```

```
RO_DATA(PAGE_SIZE)  
EXCEPTION_TABLE(8)  
NOTES  
+ KERNEL_PG_TABLES
```

The KSMA patch

- The kernel assumes the relative offset between the `swapper_pg_dir` and the `tramp_pg_dir`

arch/arm64/kernel/entry.S

```
.macro tramp_map_kernel, tmp
mrs \tmp, ttbr1_el1
sub \tmp, \tmp, #(SWAPPER_DIR_SIZE + RESERVED_TTBRO_SIZE)
bic \tmp, \tmp, #USER_ASID_FLAG
msr ttbr1_el1, \tmp
[...]
.endm
```

The KSMA patch



- Update the page table through fixmap

```
void set_swapper_pgd(pgd_t *pgdp, pgd_t pgd)
{
    pgd_t *fixmap_pgd;
    spin_lock(&swapper_pgdir_lock);
    fixmap_pgd = pgd_set_fixmap(pa(pgd));
    WRITE_ONCE(*fixmap_pgd, pgd);
    pgd_clear_fixmap();
    spin_unlock(&swapper_pgdir_lock);
}
```

The SMA mitigation



- **CONFIG_SLAB_FREELIST_HARDENED^[1]**
 - ✓ Not enabled in the kernel 4.9
 - ✓ Not enabled on the Google Pixel (4.14)
- **CONFIG_INIT_ON_FREE_DEFAULT_ON^[2]**
 - ✓ Not enabled in the Android kernel

[1] https://hardenedlinux.github.io/system-security/2017/12/02/linux_kernel_4.14%E7%9A%84SLAB_FREELIST_HARDENED%E7%9A%84%E7%AE%80%E8%A6%81%E5%88%86%E6%9E%90.html

[2] <https://outflux.net/blog/archives/2019/11/14/security-things-in-linux-v5-3/>

SLAB_FREELIST_HARDENED



- **Generate a random number**

mm/slub.c

```
int kmem_cache_open(struct kmem_cache *s, unsigned long flags)
{
    s->flags = kmem_cache_flags(s->size, flags, s->name, s->ctor);
    s->reserved = 0;
#ifdef CONFIG_SLAB_FREELIST_HARDENED
    s->random = get_random_long();
#endif
}
```

SLAB_FREELIST_HARDENED



- Obfuscate the address of SLAB

mm/slub.c

```
void *freelist_ptr(const struct kmem_cache *s, void *ptr, unsigned long ptr_addr)
{
#ifdef CONFIG_SLAB_FREELIST_HARDENED
    return (void *)((unsigned long)ptr ^ s->random ^
        (unsigned long)kasan_reset_tag((void *)ptr_addr));
#else
    return ptr;
#endif
}
```

SLAB_FREELIST_HARDENED



- The obfuscation has weaknesses

kmalloc-32 freelist walk, before:

ptr	ptr_addr	stored value	random number
ffff90c22e019020	ffff90c22e019000	is 86528eb656b3b5bd	(86528eb656b3b59d)
ffff90c22e019040	ffff90c22e019020	is 86528eb656b3b5fd	(86528eb656b3b59d)
ffff90c22e019060	ffff90c22e019040	is 86528eb656b3b5bd	(86528eb656b3b59d)

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/commit/?id=1ad53d9fa3f6168ebcf48a50e08b170432da2257>



- Improve the obfuscation

mm/slub.c

```
void *freelist_ptr(const struct kmem_cache *s, void *ptr, unsigned long ptr_addr)
{
#ifdef CONFIG_SLAB_FREELIST_HARDENED
    return (void *)((unsigned long)ptr ^ s->random ^
        (unsigned long)kasan_reset_tag((void *)ptr_addr));
+       swab((unsigned long)kasan_reset_tag((void *)ptr_addr)));
#else
    return ptr;
#endif
}
```

SLAB_FREELIST_HARDENED



- Check the double free

mm/slub.c

```
void set_freepointer(struct kmem_cache *s, void *object, void *fp)
{
    unsigned long freeptr_addr = (unsigned long)object + s->offset;
#ifdef CONFIG_SLAB_FREELIST_HARDENED
    BUG_ON(object == fp); // Double free
#endif
    *(void **)freeptr_addr = freelist_ptr(s, fp, freeptr_addr);
}
```

INIT_ON_FREE_DEFAULT_ON



- Clear the data in the released SLAB

mm/slub.c

```
bool slab_free_freelist_hook(struct kmem_cache *s, void **head, void **tail)
{
    do {
        object = next;
        next = get_freepointer(s, object);
        if (slab_want_init_on_free(s)) {
            memset(object, 0, s->object_size);
            rsize = (s->flags & SLAB_RED_ZONE) ? s->red_left_pad : 0;
            memset((char *)object + s->inuse, 0, s->size - s->inuse - rsize);
        }
    } while (object != old_tail);
}
```

Contents



- 01.** Three Dark Clouds
- 02.** Thunder and Lightning
- 03.** Dispel the Clouds and See the Sun
- 04.** Conclusion

Conclusion



- **The combination of the SMA and the KSMA is powerful**
- **In order to harden the kernel, we need:**
 1. Mitigate known attacks (the ret2dir)
 2. Fill the gap between Android and the mainline (the KSMA patch)
 3. Prevent kernel fragmentation (the SMA)

Thank You

Jun Yao @ Alpha Lab
yaojun8558363@gmail.com